

LP3 Practical 5 Implement K-Nearest Neighbors algorithm on diabetes.csv dataset. Compute confusion matrix, accuracy, error rate, precision and recall on the given dataset. Dataset link : <https://www.kaggle.com/datasets/abdallamahgoub/diabetes>

```
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
import numpy as np
```

```
df=pd.read_csv("diabetes.csv")
df
```

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	Pedigree	Age	Outcome
0	6	148	72	35	0	33.6	0.627	50	1
1	1	85	66	29	0	26.6	0.351	31	0
2	8	183	64	0	0	23.3	0.672	32	1
3	1	89	66	23	94	28.1	0.167	21	0
4	0	137	40	35	168	43.1	2.288	33	1
...	...	...	...	...	...	...	...	...	...
763	10	101	76	48	180	32.9	0.171	63	0
764	2	122	70	27	0	36.8	0.340	27	0
765	5	121	72	23	112	26.2	0.245	30	0
766	1	126	60	0	0	30.1	0.349	47	1
767	1	93	70	31	0	30.4	0.315	23	0

768 rows × 9 columns

Next steps: [Generate code with df](#) [New interactive sheet](#)

```
print(df.head())
print(df.tail())
print(df.shape)
print(df.size)
```

```
Pregnancies Glucose BloodPressure SkinThickness Insulin BMI \
0           6      148          72        35       0  33.6
1           1       85          66        29       0  26.6
2           8      183          64        0       0  23.3
3           1       89          66        23    94  28.1
4           0      137          40        35    168  43.1

Pedigree Age Outcome
0     0.627  50      1
1     0.351  31      0
2     0.672  32      1
3     0.167  21      0
4     2.288  33      1
Pregnancies Glucose BloodPressure SkinThickness Insulin BMI \
763        10      101          76        48    180  32.9
764        2       122          70        27       0  36.8
765        5      121          72        23    112  26.2
766        1      126          60        0       0  30.1
767        1       93          70        31       0  30.4

Pedigree Age Outcome
763     0.171  63      0
764     0.340  27      0
765     0.245  30      0
766     0.349  47      1
767     0.315  23      0
(768, 9)
6912
```

```
print(df.info())
print(df.describe())
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 768 entries, 0 to 767
Data columns (total 9 columns):
 #   Column      Non-Null Count  Dtype  
--- 
 0   Pregnancies    768 non-null   int64  
 1   Glucose       768 non-null   int64
```

```

2   BloodPressure    768 non-null    int64
3   SkinThickness    768 non-null    int64
4   Insulin          768 non-null    int64
5   BMI              768 non-null    float64
6   Pedigree         768 non-null    float64
7   Age              768 non-null    int64
8   Outcome          768 non-null    int64
dtypes: float64(2), int64(7)
memory usage: 54.1 KB
None
<bound method NDFrame.describe of
   Pregnancies  Glucose  BloodPressure  SkinThickness  Insulin  BMI  \
0           6      148        72          35            0     33.6
1           1       85        66          29            0     26.6
2           8      183        64          0             0     23.3
3           1       89        66          23            94    28.1
4           0      137        40          35            168    43.1
..          ...
763          10      101        76          48            180    32.9
764          2       122        70          27            0     36.8
765          5       121        72          23            112    26.2
766          1       126        60          0             0     30.1
767          1       93        70          31            0     30.4

   Pedigree  Age  Outcome
0      0.627  50      1
1      0.351  31      0
2      0.672  32      1
3      0.167  21      0
4      2.288  33      1
..          ...
763     0.171  63      0
764     0.340  27      0
765     0.245  30      0
766     0.349  47      1
767     0.315  23      0

```

[768 rows x 9 columns]>

```

# Split features and target
X = df.drop("Outcome", axis=1)    # Features
y = df["Outcome"]                 # Target

```

```

# Train-test split

from sklearn.model_selection import train_test_split
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)

# Feature scaling (important for KNN)
from sklearn.preprocessing import StandardScaler
scaler = StandardScaler()
X_train = scaler.fit_transform(X_train)
X_test = scaler.transform(X_test)

```

```

# Initialize KNN model (k=5 is common, can tune with cross-validation)
from sklearn.neighbors import KNeighborsClassifier
knn = KNeighborsClassifier(n_neighbors=5)
knn.fit(X_train, y_train)

```

▼ KNeighborsClassifier ⓘ ⓘ

KNeighborsClassifier()

```

# Predictions
y_pred = knn.predict(X_test)
y_pred

array([0, 0, 0, 0, 1, 0, 0, 1, 1, 0, 0, 0, 0, 0, 0, 1, 0, 0, 1, 0, 1, 0,
       0, 0, 1, 0, 0, 0, 0, 1, 1, 1, 1, 1, 1, 1, 0, 0, 1, 0, 1, 1, 1, 0,
       0, 0, 1, 0, 0, 1, 0, 1, 0, 0, 0, 0, 0, 0, 0, 1, 1, 0, 0, 0, 1, 0, 0, 1,
       0, 1, 0, 0, 1, 0, 0, 0, 0, 1, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 1,
       0, 0, 0, 0, 0, 0, 1, 1, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 1, 0,
       0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 1, 1, 1,
       0, 0, 0, 0, 1, 1, 1, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0])

```

```

from sklearn.metrics import confusion_matrix, accuracy_score, precision_score, recall_score
# Confusion Matrix
cm = confusion_matrix(y_test, y_pred)

# Metrics
accuracy = accuracy_score(y_test, y_pred)
error_rate = 1 - accuracy
precision = precision_score(y_test, y_pred)

```

```
recall = recall_score(y_test, y_pred)

# Display results
print("Confusion Matrix:\n", cm)
print(f"Accuracy: {accuracy:.4f}")
print(f"Error Rate: {error_rate:.4f}")
print(f"Precision: {precision:.4f}")
print(f"Recall: {recall:.4f}")

Confusion Matrix:
[[79 20]
 [27 28]]
Accuracy: 0.6948
Error Rate: 0.3052
Precision: 0.5833
Recall: 0.5091
```

